
Amos House Price Predictions

ML workflow:

1. SetUp
2. Import Data
3. Explore(EDA)
4. Splitting
5. Modelling
6. Splitting

1. SetUp

We are going to import all the necessary libraries here.

```
In [1]: import sys
import logging

import pickle
import pandas as pd
import numpy as np
import math

# creating path object
from pathlib import Path

# visualization
import matplotlib
import plotly
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns

# machine learning
import sklearn

from sklearn.linear_model import LinearRegression
```

```

from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
# preprocessing
from sklearn.preprocessing import (
    OneHotEncoder,
    OrdinalEncoder,
    FunctionTransformer,
    PolynomialFeatures,
    StandardScaler
)
# imputing missing values
from sklearn.impute import SimpleImputer
# evaluation metrics
from sklearn.metrics import r2_score, mean_absolute_percentage_error
# compose
from sklearn.compose import ColumnTransformer
# making pipeline
from sklearn.pipeline import Pipeline
# feature selections
from sklearn.feature_selection import VarianceThreshold
# model selections eg. splitting
from sklearn.model_selection import train_test_split, learning_curve, GridSearchCV
# pca decomposition
from sklearn.decomposition import PCA
# evaluation metrics
from sklearn.metrics import mean_absolute_percentage_error, mean_absolute_error, r2_score

```

We have all the libraries in place. Let us print our library versions. This step ensures reproducibility

```

In [2]: # Printing version of our libraries
print("Platform: ", sys.platform)
print("Python: ", sys.version)
print("---")
print("Matplotlib: ", matplotlib.__version__)
print("Pandas: ", pd.__version__)
print("Seaborn: ", sns.__version__)
print("Plotly Express: ", plotly.__version__)
print("Numpy: ", np.__version__)
print("Sklearn: ", sklearn.__version__)

```

```
Platform: win32
Python: 3.13.2 (tags/v3.13.2:4f8bb39, Feb 4 2025, 15:23:48) [MSC v.1942 64 bit (AMD64)]
---
Matplotlib: 3.10.0
Pandas: 2.2.3
Seaborn: 0.13.2
Plotly Express: 6.0.0
Numpy: 2.2.2
Sklearn: 1.6.1
```

Define the logging configurations.

```
In [3]: # Configure
config_path = Path.cwd()/"Training"/"Configure"
config_path.mkdir(parents=True, exist_ok=True)
logging.basicConfig(
    level=logging.INFO,
    filename = config_path / "logging.log",
)
```

Let us not define matplotlib configurations.

```
In [4]: #Matplotlib configuration
plt.rc("font", size=12)
plt.rc("axes", labelsz=12, titlesz=14)
plt.rc("legend", fontsize=8)
plt.rc("xtick", labelsz=10)
plt.rc("ytick", labelsz=10)
%matplotlib inline
```

Making two functions that will help us saving images and the other saving a trained model.

```
In [5]: # save figure function
def save_plot(fname, filetype, fig=None, dpi=300, tight_layout=True, format="png"):
    """Saving the plot as an image
    Save a plot image within ``Train/Images`` folders. The fname will be the name of the image with the
    extension .(format).``Note`` the default format is png.
    The plot by default will be saved under 300 resolution as inches.

    Parameters:
    -----
    fname: str
        ->String object for the name of the plot.
    filetype: str
        -> Plot type eg plt(matplotlib), or px(plotly.express)
    fig: px.Figure
```

```

-> figure object form plotly.express
dpi: int
    -> Numerical variable for the pixel resolution.
tight_layout: bool
    -> If true the plot will be save on a tight layout.
format: str
    -> String object for the image extension. By default is 'png' but we can have: 'jpeg', 'jpg', etc..

```

Returns:

None

"""

```

# Root path
image_path = Path.cwd() / "Training" / "Images"
# Making the folders
logging.info("Creting image path")
image_path.mkdir(parents=True, exist_ok=True)
# Image name
image_name = image_path / f"{fname}.{format}"
# Layout format
if filetype == "plt":
    if tight_layout:
        plt.tight_layout()
    # Saving the plot
    logging.info(f"Saving the plot as {fname}.{format}")
    plt.savefig(fname=image_name, dpi=dpi, format=format)
    # logging saving
    logging.info(f"Sucess! Saved the plot as {fname}.{format}")
elif filetype == "px":
    logging.info(f"Saving the plot as {fname}.{format}")
    # writting the image
    fig.write_image(file= image_name, format= format)
    logging.info(f"Sucess! Saved the plot as {fname}.{format}")

```

```

In [6]: # Saving the model
def save_model(mname, model):
    """Saving the model.
    Get the model and save it using pickle. The model will have the name from mname.

```

Parameters:

mname: str

 Name of the model as a string object

model: sklearn.model

 The trained model

Returns:

```

None
"""

# Root path
model_path = Path.cwd() / "Training" / "Models"
# Making the folders
logging.info("Creating Model path")
model_path.mkdir(parents=True, exist_ok=True)
# Model name
model_name = model_path / f"{mname}.pkl"
logging.info(f"Saving the model as {mname}.pkl")
# Creating the pickle file
with open(model_name, "wb") as f:
    pickle.dump(model, f)
# final log
logging.info(f"Success! Saved the model as {mname}.pkl")

```

We have a solid setup sections, let us start data importation.

2. Import and EDA

I have created a Training module where I have all my classes. We are going to get the `WrangleRepository` class that does the following:

1. Get the data from the csv file
2. Do a basic cleaning
3. Feature selection
4. Feature engineering
5. Outlier removing

In [7]: `from Training import WrangleRepository`

```

# instantiating the class
repo = WrangleRepository()
print(type(repo))
repo

```

```
<class 'Training.WrangleRepository'>
```

Out[7]: `WrangleRepository filepath=C:\Users\MY PC\Desktop\Projects\Regression\AmosHousePriceModelling\train.csv`

Let us now use the function to load the data in `Desktop/Projects/Regression/AmosHousePriceModelling/train.csv`

In [8]: `repo.wrangle()`
`df = repo.get_data("wrangled")`

```
print(df.shape)
df.head()
```

(1460, 80)

Out[8]:

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | LandContour | Utilities | LotConfig | ... | PoolArea | PoolQC | Fence | ... |
|----|------------|----------|-------------|---------|--------|-------|----------|-------------|-----------|-----------|-----|----------|--------|-------|-----|
| Id | | | | | | | | | | | | | | | |
| 1 | 60 | RL | 65.0 | 8450 | Pave | NaN | Reg | Lvl | AllPub | Inside | ... | 0 | NaN | NaN | |
| 2 | 20 | RL | 80.0 | 9600 | Pave | NaN | Reg | Lvl | AllPub | FR2 | ... | 0 | NaN | NaN | |
| 3 | 60 | RL | 68.0 | 11250 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | NaN | NaN | |
| 4 | 70 | RL | 60.0 | 9550 | Pave | NaN | IR1 | Lvl | AllPub | Corner | ... | 0 | NaN | NaN | |
| 5 | 60 | RL | 84.0 | 14260 | Pave | NaN | IR1 | Lvl | AllPub | FR2 | ... | 0 | NaN | NaN | |

5 rows × 80 columns



We have our data successfully. We will start by doing basic data cleaning.

In [9]: `df.describe().T`

Out[9]:

| | count | mean | std | min | 25% | 50% | 75% | max |
|---------------------|--------|--------------|-------------|--------|---------|--------|----------|----------|
| MSSubClass | 1460.0 | 56.897260 | 42.300571 | 20.0 | 20.00 | 50.0 | 70.00 | 190.0 |
| LotFrontage | 1201.0 | 70.049958 | 24.284752 | 21.0 | 59.00 | 69.0 | 80.00 | 313.0 |
| LotArea | 1460.0 | 10516.828082 | 9981.264932 | 1300.0 | 7553.50 | 9478.5 | 11601.50 | 215245.0 |
| OverallQual | 1460.0 | 6.099315 | 1.382997 | 1.0 | 5.00 | 6.0 | 7.00 | 10.0 |
| OverallCond | 1460.0 | 5.575342 | 1.112799 | 1.0 | 5.00 | 5.0 | 6.00 | 9.0 |
| YearBuilt | 1460.0 | 1971.267808 | 30.202904 | 1872.0 | 1954.00 | 1973.0 | 2000.00 | 2010.0 |
| YearRemodAdd | 1460.0 | 1984.865753 | 20.645407 | 1950.0 | 1967.00 | 1994.0 | 2004.00 | 2010.0 |
| MasVnrArea | 1452.0 | 103.685262 | 181.066207 | 0.0 | 0.00 | 0.0 | 166.00 | 1600.0 |
| BsmtFinSF1 | 1460.0 | 443.639726 | 456.098091 | 0.0 | 0.00 | 383.5 | 712.25 | 5644.0 |
| BsmtFinSF2 | 1460.0 | 46.549315 | 161.319273 | 0.0 | 0.00 | 0.0 | 0.00 | 1474.0 |
| BsmtUnfSF | 1460.0 | 567.240411 | 441.866955 | 0.0 | 223.00 | 477.5 | 808.00 | 2336.0 |
| TotalBsmtSF | 1460.0 | 1057.429452 | 438.705324 | 0.0 | 795.75 | 991.5 | 1298.25 | 6110.0 |
| 1stFlrSF | 1460.0 | 1162.626712 | 386.587738 | 334.0 | 882.00 | 1087.0 | 1391.25 | 4692.0 |
| 2ndFlrSF | 1460.0 | 346.992466 | 436.528436 | 0.0 | 0.00 | 0.0 | 728.00 | 2065.0 |
| LowQualFinSF | 1460.0 | 5.844521 | 48.623081 | 0.0 | 0.00 | 0.0 | 0.00 | 572.0 |
| GrLivArea | 1460.0 | 1515.463699 | 525.480383 | 334.0 | 1129.50 | 1464.0 | 1776.75 | 5642.0 |
| BsmtFullBath | 1460.0 | 0.425342 | 0.518911 | 0.0 | 0.00 | 0.0 | 1.00 | 3.0 |
| BsmtHalfBath | 1460.0 | 0.057534 | 0.238753 | 0.0 | 0.00 | 0.0 | 0.00 | 2.0 |
| FullBath | 1460.0 | 1.565068 | 0.550916 | 0.0 | 1.00 | 2.0 | 2.00 | 3.0 |
| HalfBath | 1460.0 | 0.382877 | 0.502885 | 0.0 | 0.00 | 0.0 | 1.00 | 2.0 |
| BedroomAbvGr | 1460.0 | 2.866438 | 0.815778 | 0.0 | 2.00 | 3.0 | 3.00 | 8.0 |
| KitchenAbvGr | 1460.0 | 1.046575 | 0.220338 | 0.0 | 1.00 | 1.0 | 1.00 | 3.0 |
| TotRmsAbvGrd | 1460.0 | 6.517808 | 1.625393 | 2.0 | 5.00 | 6.0 | 7.00 | 14.0 |
| Fireplaces | 1460.0 | 0.613014 | 0.644666 | 0.0 | 0.00 | 1.0 | 1.00 | 3.0 |
| GarageYrBlt | 1379.0 | 1978.506164 | 24.689725 | 1900.0 | 1961.00 | 1980.0 | 2002.00 | 2010.0 |
| GarageCars | 1460.0 | 1.767123 | 0.747315 | 0.0 | 1.00 | 2.0 | 2.00 | 4.0 |

| | count | mean | std | min | 25% | 50% | 75% | max |
|----------------------|--------|---------------|--------------|---------|-----------|----------|-----------|----------|
| GarageArea | 1460.0 | 472.980137 | 213.804841 | 0.0 | 334.50 | 480.0 | 576.00 | 1418.0 |
| WoodDeckSF | 1460.0 | 94.244521 | 125.338794 | 0.0 | 0.00 | 0.0 | 168.00 | 857.0 |
| OpenPorchSF | 1460.0 | 46.660274 | 66.256028 | 0.0 | 0.00 | 25.0 | 68.00 | 547.0 |
| EnclosedPorch | 1460.0 | 21.954110 | 61.119149 | 0.0 | 0.00 | 0.0 | 0.00 | 552.0 |
| 3SsnPorch | 1460.0 | 3.409589 | 29.317331 | 0.0 | 0.00 | 0.0 | 0.00 | 508.0 |
| ScreenPorch | 1460.0 | 15.060959 | 55.757415 | 0.0 | 0.00 | 0.0 | 0.00 | 480.0 |
| PoolArea | 1460.0 | 2.758904 | 40.177307 | 0.0 | 0.00 | 0.0 | 0.00 | 738.0 |
| MiscVal | 1460.0 | 43.489041 | 496.123024 | 0.0 | 0.00 | 0.0 | 0.00 | 15500.0 |
| MoSold | 1460.0 | 6.321918 | 2.703626 | 1.0 | 5.00 | 6.0 | 8.00 | 12.0 |
| YrSold | 1460.0 | 2007.815753 | 1.328095 | 2006.0 | 2007.00 | 2008.0 | 2009.00 | 2010.0 |
| SalePrice | 1460.0 | 180921.195890 | 79442.502883 | 34900.0 | 129975.00 | 163000.0 | 214000.00 | 755000.0 |

As seen, we have 1460 number of houses. Some of the features have missing values and we need to work on that. Also it is important to note that some features have outliers for example `lotArea`. This is determined from the sense that, we have a mean of approximately 10,500 and std of 9900 with the maximum value being at 255,000. Clearly we have outlier and most likely our data will be skewed.


```
In [10]: # checking for missing values
missing_values = (df.isnull().sum()[df.isnull().sum() > 1]).sort_values()
print(f"Missing values in our Features: \n{missing_values}")
```


Missing values in our Features:

| | |
|--------------|------|
| MasVnrArea | 8 |
| BsmtQual | 37 |
| BsmtCond | 37 |
| BsmtFinType1 | 37 |
| BsmtFinType2 | 38 |
| BsmtExposure | 38 |
| GarageCond | 81 |
| GarageQual | 81 |
| GarageFinish | 81 |
| GarageYrBlt | 81 |
| GarageType | 81 |
| LotFrontage | 259 |
| FireplaceQu | 690 |
| MasVnrType | 872 |
| Fence | 1179 |
| Alley | 1369 |
| MiscFeature | 1406 |
| PoolQC | 1453 |

dtype: int64


```
In [11]: # plotting the top 10 most missing values
missing_values.tail(10).plot(kind="barh")
plt.xlabel("Missing Values (counts)")
plt.ylabel("Independent Variables (features)")
plt.title("Top Features: With Missing Values")
plt.grid(linestyle="--", axis="x")
save_plot(fname="top10_missing_values", filetype="plt")
plt.show()
```



```
In [12]: # percentage counts
missing_values_pct = pd.Series(((100 * missing_values.values / len(df))),
                                index=missing_values.index, name="missing_pct")
missing_values_pct
```

```
Out[12]: MasVnrArea      0.547945
BsmtQual      2.534247
BsmtCond      2.534247
BsmtFinType1  2.534247
BsmtFinType2  2.602740
BsmtExposure  2.602740
GarageCond     5.547945
GarageQual     5.547945
GarageFinish   5.547945
GarageYrBlt    5.547945
GarageType     5.547945
LotFrontage   17.739726
FireplaceQu   47.260274
MasVnrType    59.726027
Fence         80.753425
Alley         93.767123
MiscFeature   96.301370
PoolQC        99.520548
Name: missing_pct, dtype: float64
```

```
In [13]: # plotting the top 10 pct features with missing values
missing_values_pct.tail(10).plot(kind="barh")
plt.xlabel("Missing Values (%)")
plt.ylabel("Independent Variables ")
plt.title("Top Features: With Missing Values in %")
plt.grid(linestyle="--", axis="x")
save_plot(fname="top10_missing_values_pct", filetype="plt")
plt.show()
```



The above plot tells us that some features like pool are not present in many properties. We need to remove those features which have many missing values. To do that we will use VarianceThreshold object to compute those features which will be useless to our model.

Some advantages of removing low variance features are:

1. Reduce features dimensionality
2. Improve model performance by reducing overfitting
3. Reduce training time

But before we do Variance reduction we are going to drop those feature with over 50% missing values.

```
In [14]: # getting columns with over 60% missing values
mask = missing_values_pct > 50
missing_cols = missing_values_pct[mask].index.to_list()
print(type(missing_cols))
print(missing_cols)
```

```
<class 'list'>
['MasVnrType', 'Fence', 'Alley', 'MiscFeature', 'PoolQC']
```

```
In [15]: repo.basic_cleaning(missing_values_pct=missing_values_pct)
df = repo.get_data("basic")
print(df.shape)
df.head()
```

(1460, 75)

```
Out[15]:
```

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | LotShape | LandContour | Utilities | LotConfig | LandSlope | ... | EnclosedPorch | 3SsnPoi |
|----|------------|----------|-------------|---------|--------|----------|-------------|-----------|-----------|-----------|-----|---------------|---------|
| Id | | | | | | | | | | | | | |
| 1 | 60 | RL | 65.0 | 8450 | Pave | Reg | Lvl | AllPub | Inside | Gtl | ... | 0 | |
| 2 | 20 | RL | 80.0 | 9600 | Pave | Reg | Lvl | AllPub | FR2 | Gtl | ... | 0 | |
| 3 | 60 | RL | 68.0 | 11250 | Pave | IR1 | Lvl | AllPub | Inside | Gtl | ... | 0 | |
| 4 | 70 | RL | 60.0 | 9550 | Pave | IR1 | Lvl | AllPub | Corner | Gtl | ... | 272 | |
| 5 | 60 | RL | 84.0 | 14260 | Pave | IR1 | Lvl | AllPub | FR2 | Gtl | ... | 0 | |

5 rows × 75 columns



Done, we have dropped those features with so many missing values. Let us go ahead and check the variance for all the numerical features first.

```
In [16]: repo.feature_selection()
df = repo.get_data("selected")
print(df.shape)
df.head()
```

(1460, 67)

Out[16]:

| | MSSubClass | LotFrontage | LotArea | OverallQual | OverallCond | YearBuilt | YearRemodAdd | MasVnrArea | BsmtFinSF1 | BsmtFinSF2 | ... | Centra |
|----|------------|-------------|---------|-------------|-------------|-----------|--------------|------------|------------|------------|-----|--------|
| Id | | | | | | | | | | | | |
| 1 | 60 | 65.0 | 8450 | 7 | 5 | 2003 | 2003 | 196.0 | 706 | 0 | ... | |
| 2 | 20 | 80.0 | 9600 | 6 | 8 | 1976 | 1976 | 0.0 | 978 | 0 | ... | |
| 3 | 60 | 68.0 | 11250 | 7 | 5 | 2001 | 2002 | 162.0 | 486 | 0 | ... | |
| 4 | 70 | 60.0 | 9550 | 7 | 5 | 1915 | 1970 | 0.0 | 216 | 0 | ... | |
| 5 | 60 | 84.0 | 14260 | 8 | 5 | 2000 | 2000 | 350.0 | 655 | 0 | ... | |

5 rows × 67 columns



Let us check low and high cardinal features.

In [17]:


```
df.select_dtypes(include="object").nunique()
```

```
Out[17]: MSZoning      5
         LotShape      4
         LandContour    4
         LotConfig      5
         LandSlope      3
         Neighborhood   25
         Condition1     9
         BldgType       5
         HouseStyle     8
         RoofStyle      6
         Exterior1st    15
         Exterior2nd    16
         ExterQual      4
         ExterCond      5
         Foundation     6
         BsmtQual       4
         BsmtCond       4
         BsmtExposure   4
         BsmtFinType1   6
         BsmtFinType2   6
         HeatingQC      5
         CentralAir     2
         Electrical     5
         KitchenQual    4
         Functional     7
         FireplaceQu    5
         GarageType     6
         GarageFinish   3
         PavedDrive     3
         SaleType       9
         SaleCondition   6
         dtype: int64
```

I think there are no cardinal features. Let us check how our features are correlated with one another.

```
In [18]: corr = df.select_dtypes(include="number").drop(columns="SalePrice").corr()
         sns.heatmap(corr, cmap="coolwarm")
         plt.title("Features Multicollinearity")
         save_plot(fname="multicollinearity", filetype="plt")
```

Features Multicollinearity



```
In [19]: # which features have a correlation above 90%
corr_matrix = abs(df.select_dtypes(include="number").drop(columns="SalePrice").corr())
upper_triangle = np.triu(np.ones(corr_matrix.shape), k=1).astype(bool)
upper_matrix = corr_matrix.where(upper_triangle)
upper_matrix
```


Out[19]:

[illegible]

| | MSSubClass | LotFrontage | LotArea | OverallQual | OverallCond | YearBuilt | YearRemodAdd | MasVnrArea | BsmtFinSF1 | BsmtFinSF |
|---------------|------------|-------------|---------|-------------|-------------|-----------|--------------|------------|------------|-----------|
| WoodDeckSF | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| OpenPorchSF | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| EnclosedPorch | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| 3SsnPorch | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| ScreenPorch | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| PoolArea | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| MiscVal | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| MoSold | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |
| YrSold | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Na |

35 rows × 35 columns

```
In [20]: high_corr = [col for col in upper_matrix.columns
                      if any(upper_matrix[col] > 0.9)
                      ]
high_corr
```

Out[20]: []

So we do have highly correlated features. Let us go ahead now and do feature engineering.

```
In [21]: sub_class = {
          20: "1-STORY 1946 & NEWER ALL STYLES",
          30: "1-STORY 1945 & OLDER",
          40: "1-STORY W/FINISHED ATTIC ALL AGES",
          45: "1-1/2 STORY - UNFINISHED ALL AGES",
          50: "1-1/2 STORY FINISHED ALL AGES",
          60: "2-STORY 1946 & NEWER",
          70: "2-STORY 1945 & OLDER",
          75: "2-1/2 STORY ALL AGES",
          80: "SPLIT OR MULTI-LEVEL",
          85: "SPLIT FOYER",
          90: "DUPLEX - ALL STYLES AND AGES",
          120: "1-STORY PUD (Planned Unit Development) - 1946 & NEWER",
          150: "1-1/2 STORY PUD - ALL AGES",
          160: "2-STORY PUD - 1946 & NEWER",
          180: "PUD - MULTILEVEL - INCL SPLIT LEV/FOYER",
```

```
190:         "2 FAMILY CONVERSION - ALL STYLES AND AGES"
    }

print(type(sub_class))
print(sub_class)
```

<class 'dict'>
{20: '1-STORY 1946 & NEWER ALL STYLES', 30: '1-STORY 1945 & OLDER', 40: '1-STORY W/FINISHED ATTIC ALL AGES', 45: '1-1/2 STORY - UNFINISHED ALL AGES', 50: '1-1/2 STORY FINISHED ALL AGES', 60: '2-STORY 1946 & NEWER', 70: '2-STORY 1945 & OLDER', 75: '2-1/2 STORY ALL AGES', 80: 'SPLIT OR MULTI-LEVEL', 85: 'SPLIT FOYER', 90: 'DUPLEX - ALL STYLES AND AGES', 120: '1-STORY PUD (Planned Unit Development) - 1946 & NEWER', 150: '1-1/2 STORY PUD - ALL AGES', 160: '2-STORY PUD - 1946 & NEWER', 180: 'PUD - MULTILEVEL - INCL SPLIT LEVEL/FOYER', 190: '2 FAMILY CONVERSION - ALL STYLES AND AGES'}

```
In [22]: repo.feature_engineering(sub_class=sub_class)
df = repo.get_data("engineered")
print(df.shape)
df.head()
```

(1460, 73)

Out[22]:

| | MSSubClass | LotFrontage | LotArea | OverallQual | OverallCond | YearBuilt | YearRemodAdd | MasVnrArea | BsmtFinSF1 | BsmtFinSF2 | ... | Garage |
|--|------------|-------------|---------|-------------|-------------|-----------|--------------|------------|------------|------------|-----|--------|
|--|------------|-------------|---------|-------------|-------------|-----------|--------------|------------|------------|------------|-----|--------|

| Id | | | | | | | | | | | | |
|----|--|------|-------|---|---|------|------|-------|-----|---|-----|--|
| 1 | 2-STORY 1946 & NEWER | 65.0 | 8450 | 7 | 5 | 2003 | 2003 | 196.0 | 706 | 0 | ... | |
| 2 | 1-STORY 1946 & NEWER ALL STYLES | 80.0 | 9600 | 6 | 8 | 1976 | 1976 | 0.0 | 978 | 0 | ... | |
| 3 | 2-STORY 1946 & NEWER | 68.0 | 11250 | 7 | 5 | 2001 | 2002 | 162.0 | 486 | 0 | ... | |
| 4 | 2-STORY 1945 & OLDER | 60.0 | 9550 | 7 | 5 | 1915 | 1970 | 0.0 | 216 | 0 | ... | |
| 5 | 2-STORY 1946 & NEWER | 84.0 | 14260 | 8 | 5 | 2000 | 2000 | 350.0 | 655 | 0 | ... | |

5 rows × 73 columns



Note we still have missing values and outliers. In the case of missing value, we will impute them inside a pipeline.

Now let us first check the distributions of our features. We will do this step manually step by step. In We will be referencing the describe objects for our dataframe.

Features with most outliers:

1. LotArea:

- Mean: 10,500
- std: 9980
- max: 215,000
- Loss: $\frac{215,000}{10,500} = 20.5$

2. LowQualFinSF:

- Mean: 6
- std: 48.6
- max: 572
- Loss: $\frac{572}{49} = 11.6$


3. MiscVal:

- Mean: 43.5
- std: 496
- Max: 15,500
- Loss: $\frac{15,500}{496} = 31.6$


We will look into those three features and see those outliers and remove them.

```
In [23]: # columns to check/huge outliers
cols = ["LotArea", "LowQualFinSF", "MiscVal"]
for col in cols:
    plt.figure(figsize=(8, 2))
    sns.boxplot(x=df[col])
    plt.title(f"Distribution: {col} Boxplot")
    save_plot(fname=f"Distribution_outlier_{col}", filetype="plt")
    plt.show()
```


Distribution: LotArea Boxplot



Distribution: LowQualFinSF Boxplot



Distribution: MiscVal Boxplot



Indeed we have outliers in our columns and we need to remove them. We will remove outlier in the upper quantile(90%)

```
In [24]: # removed outlier
repo.remove_outliers()
df = repo.get_data()
print(df.shape)
df.head()
```

(1170, 73)

Out[24]: MSSubClass LotFrontage LotArea OverallQual OverallCond YearBuilt YearRemodAdd MasVnrArea BsmtFinSF1 BsmtFinSF2 ... Garage

| Id | | | | | | | | | | | |
|----|--|------|-------|---|---|------|------|-------|-----|---|-----|
| 1 | 2-STORY 1946 & NEWER | 65.0 | 8450 | 7 | 5 | 2003 | 2003 | 196.0 | 706 | 0 | ... |
| 2 | 1-STORY 1946 & NEWER ALL STYLES | 80.0 | 9600 | 6 | 8 | 1976 | 1976 | 0.0 | 978 | 0 | ... |
| 3 | 2-STORY 1946 & NEWER | 68.0 | 11250 | 7 | 5 | 2001 | 2002 | 162.0 | 486 | 0 | ... |
| 5 | 2-STORY 1946 & NEWER | 84.0 | 14260 | 8 | 5 | 2000 | 2000 | 350.0 | 655 | 0 | ... |
| 6 | 1-1/2 STORY FINISHED ALL AGES | 85.0 | 14115 | 5 | 5 | 1993 | 1995 | 0.0 | 732 | 0 | ... |

5 rows × 73 columns



4. Splitting

Now that we have our class well defined and we also have the most current data, the next thing we want to do is split our data into train and validation set.

```
In [25]: # Vertical split
target = "SalePrice"
X = df.drop(columns=target)
y = df[target]
print(f"X shape: {X.shape}")
print(f"y shape: {y.shape}")
```

X shape: (1170, 72)
y shape: (1170,)

```
In [26]: # Horizontal split
X_train, X_val, y_train, y_val = train_test_split(
```

```

X, y, test_size=0.15
)
print(f"X_train shape: {X_train.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"X_val shape: {X_val.shape}")
print(f"y_val shape: {y_val.shape}")

```

```

X_train shape: (994, 72)
y_train shape: (994,)
X_val shape: (176, 72)
y_val shape: (176,)

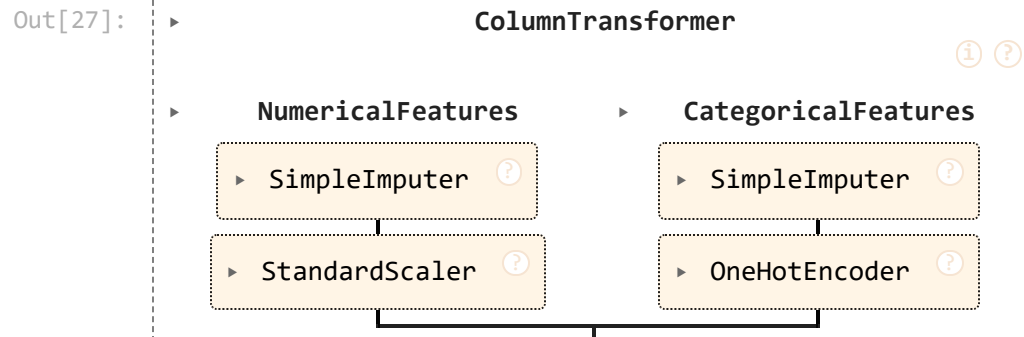
```

5. Pipeline

```

In [27]: # numerical pipeline
num_pipeline = Pipeline([
    ("imputer", SimpleImputer(strategy="mean")),
    ("scaler", StandardScaler())
])
# categorical pipeline
cat_pipeline = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(handle_unknown="ignore"))
])
# column transformer
col_pipeline = ColumnTransformer([
    ("NumericalFeatures", num_pipeline, X_train.select_dtypes(include="number").columns),
    ("CategoricalFeatures", cat_pipeline, X_train.select_dtypes(include="object").columns)
])
col_pipeline

```



6. PCA Decomposition

In order to visualize our feature matrix (independent variables) against dependent variable `SalePrice`, we need to reduce the number of dimensions of the feature matrix. In our case we want just one dimension.

We will continue to build the pipeline and created a `pca_pipeline` that we will then fit and transform into a single vector, of 1-dimension matrix.

```
In [28]: # building the pipeline
pca_pipeline = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("PCA Algorithm", PCA(n_components=1, random_state=42))
    ]
)
# fitting and transforming
X_t = pca_pipeline.fit_transform(X_train)
print("Type of X_t: ", type(X_t))
print("Total Number of items: ", len(X_t))
print("Number of dimensions: ", X_t.ndim)
print("Transformed X_train: \n", X_t[:4])
```

Type of X_t: <class 'numpy.ndarray'>

Total Number of items: 994

Number of dimensions: 2

Transformed X_train:

```
[[ -1.86861579]
 [ -3.87269582]
 [  0.26245341]
 [  2.75529971]]
```

Now that we have that, the next thing we want to do is making a function that will help us make the image instantly.

```
In [29]: def scatter_plot(x, y, y_pred, label):
        """
        Make a scatter plot comparing actual vs. predicted values.

        Parameters:
            x: array-like
                Independent variable (e.g., PCA-transformed features)
            y: array-like
                Actual target values (e.g., true Sale Prices)
            y_pred: array-like
                Predicted target values
            label: str
                Label for the plot title
        """
        df = pd.DataFrame({
            "x": x.ravel(),
```



```

        "y": y,
        "y_pred": y_pred
    })

    df_melt = pd.melt(
        frame=df,
        id_vars="x",
        value_vars=["y", "y_pred"],
        var_name="Set",
        value_name="Sale Price"
    )

    fig = px.scatter(
        data_frame=df_melt,
        x="x",
        y="Sale Price",
        color="Set",
        title=f"{label} Scatter Plot: Decomposed Features vs. Sale Price"
    )

    fig.update_layout(
        xaxis_title="Decomposed Feature(s)",
        yaxis_title="Sale Price ($)",
        legend_title="Plot Type",
        template = "plotly_white"
    )

    # return
    return fig

```

```

In [30]: fig = scatter_plot(X_t, y_train, y_pred=None, label="")
         save_plot(fname="DecomposedScatter", filetype="plt")
         fig.show()

```

<Figure size 640x480 with 0 Axes>

We have our beautiful scatter plot and it seems like our features follow a polynomial kind of a function. We should keep this in mind. But for now let us have our baseline model.

7. Baseline and Linear Regression Model

```
In [31]: y_mean = y_train.mean()
baseline_model = len(y_train) * [y_mean]
print("Mean Sale Price: ", int(y_mean))
```

Mean Sale Price: 175846

We have our baseline model, let us evaluate it's performance using mean absolute error which is computed as

$$MAE = \frac{1}{n} \sum_i^n (y_i - \hat{y}_i)$$

$$MAPE = \frac{100}{n} \sum_i^n \frac{(y_i - \hat{y}_i)^2}{y_i}$$

$$\text{And finally computing coefficient of determination } r^2 \text{ score as: } R^2 = 1 - \frac{\sum_i^n (y_i - \hat{y})^2}{\sum_i^n (y_i - \bar{y})^2}$$

Note:

1. y_i - Actual dependent values
2. \hat{y}_i - Predicted values
3. \bar{y} - Mean
4. n - Number of samples

```
In [32]: mae = mean_absolute_error(y_train, baseline_model)
mape = mean_absolute_percentage_error(y_train, baseline_model)
cod = r2_score(y_train, baseline_model)
print(f"MAE: ${np.round(mae, 2)}")
print(f"MAPE: {100 * np.round(mape, 2)}%")
print(f"R2: {100 * np.round(cod, 2)}%")
```

MAE: \$51164.6

MAPE: 33.0%

R2: 0.0%

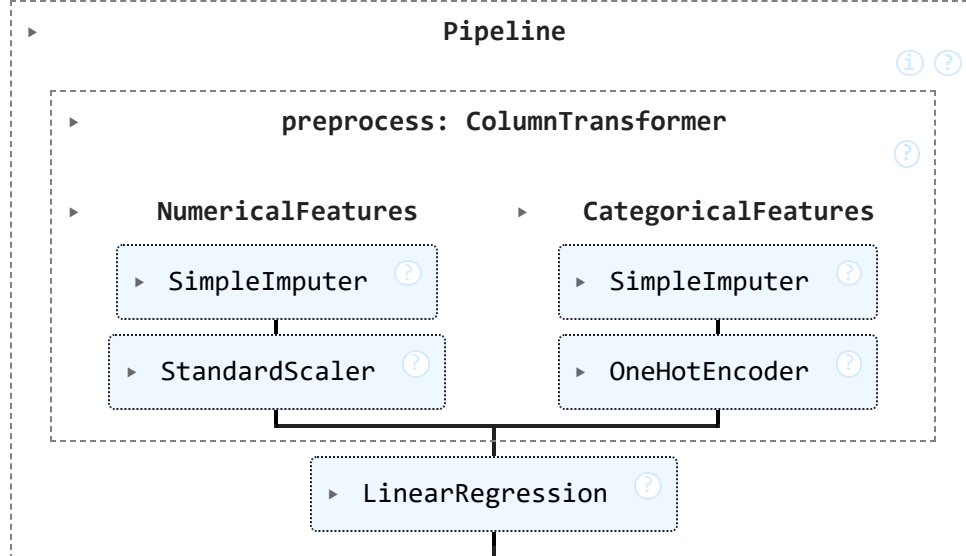
```
In [33]: fig = scatter_plot(X_t, y_train, y_pred=baseline_model, label="Baseline Model")
save_plot(fname="BaselineModelScatter", filetype="plt")
fig.show()
```

<Figure size 640x480 with 0 Axes>

We are off by 36% from the actual values. This is a reasonable range but there is room for improvement. We will train a linear regression model first.

```
In [34]: # Linear regression
linear_model = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("linear_model", LinearRegression())
    ]
)
# Training the model
linear_model.fit(X_train, y_train)
```

Out[34]:



```
In [35]: # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = linear_model.predict(X_train)
# With training data
mae = mean_absolute_error(y_train, pred)
mape = mean_absolute_percentage_error(y_train, pred)
cod = r2_score(y_train, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$11391.7

Training MAPE: 7.000000000000001%

Training R2: 94.0%

We are off by 7% with our true values. Next let us evaluate our model with the validation dataset. We will check if there is any form of overfitting that needs to be addressed.

```
In [109... # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = linear_model.predict(X_val)
# With training data
mae = mean_absolute_error(y_val, pred)
mape = mean_absolute_percentage_error(y_val, pred)
cod = r2_score(y_val, pred)
print(f"Validation MAE: ${np.round(mae, 2)}")
print(f"Validation MAPE: {100 * np.round(mape, 2)}%")
print(f"Validation R2: {100 * np.round(cod, 2)}%")
```

Validation MAE: \$16265.29
Validation MAPE: 9.0%
Validation R2: 88.0%

In [37]: `from Training import LearningCurve`

```
lc = LearningCurve(estimator=linear_model, X=X, y=y)
print(type(lc))
lc
```

<class 'Training.LearningCurve'>

Out[37]: LearningCurve: <class 'sklearn.pipeline.Pipeline'>

In [38]: `# Getting learning curve list`

```
lc.learning_curve()
data = lc.get_data("lc")
print(type(data))
data
```

[learning_curve] Training set sizes: [93 304 514 725 936]

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 25 out of 25 | elapsed: 10.5s finished

<class 'list'>

Out[38]: `array([93, 304, 514, 725, 936]),`
`array([1. , 0.97686676, 0.95751753, 0.94935131, 0.9460286]),`
`array([0.43860028, 0.71470291, 0.85816672, 0.87820229, 0.87343716])]`

In [39]: `# Making the dataframe`

```
lc.make_dataframe()
df_lc = lc.get_data("df_lc")
print(type(df_lc))
df_lc.head()
```

<class 'pandas.core.frame.DataFrame'>

Out[39]:

| | Train Size | Train R2 | Validation R2 |
|--|------------|----------|---------------|
|--|------------|----------|---------------|

| | | | |
|---|-----|----------|----------|
| 0 | 93 | 1.000000 | 0.438600 |
| 1 | 304 | 0.976867 | 0.714703 |
| 2 | 514 | 0.957518 | 0.858167 |
| 3 | 725 | 0.949351 | 0.878202 |
| 4 | 936 | 0.946029 | 0.873437 |

```
In [40]: # melt our dataframe
lc.melt_dataframe()
df_melt = lc.get_data("melt_lc")
print(type(df_melt))
df_melt.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Out[40]:
```

| | Train Size | Set | R2 |
|----------|-------------------|------------|-----------|
| 0 | 93 | Train R2 | 1.000000 |
| 1 | 304 | Train R2 | 0.976867 |
| 2 | 514 | Train R2 | 0.957518 |
| 3 | 725 | Train R2 | 0.949351 |
| 4 | 936 | Train R2 | 0.946029 |

```
In [41]: # plotting the figure
lc.plot_lc()
fig = lc.get_data()
fig.show()
```

With the above information we might want to build a class called LearningCurve that will have all that information.

That is good information. Next thing we want to do is plot the scatter plot and seen.

```
In [42]: val_pred = linear_model.predict(X_val)
val_X_t = pca_pipeline.fit_transform(X_val)
```

```
In [43]: fig = scatter_plot(x=val_X_t, y=y_val, y_pred=val_pred, label="Linear Model")
save_plot(fname="LinerModelScatter", filetype="plt")
fig.show()
```


<Figure size 640x480 with 0 Axes>

```
In [44]: # Let us save the model  
save_model(mname="LinearRegressionModel", model=linear_model)
```

8. ID mapping

```
In [45]: repo = WrangleRepository(file_name="test (1).csv")  
repo
```

```
Out[45]: WrangleRepository filepath=C:\Users\MY PC\Desktop\Projects\Regression\AmosHousePriceModelling\test (1).csv
```

```
In [46]: # getting the data  
repo.wrangle()
```

```
df = repo.get_data("wrangled")
print(df.shape)
df.head()
```

(1459, 79)

Out[46]:

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | LandContour | Utilities | LotConfig | ... | ScreenPorch | PoolArea | F |
|------|------------|----------|-------------|---------|--------|-------|----------|-------------|-----------|-----------|-----|-------------|----------|---|
| Id | | | | | | | | | | | | | | |
| 1461 | 20 | RH | 80.0 | 11622 | Pave | NaN | Reg | Lvl | AllPub | Inside | ... | 120 | 0 | |
| 1462 | 20 | RL | 81.0 | 14267 | Pave | NaN | IR1 | Lvl | AllPub | Corner | ... | 0 | 0 | |
| 1463 | 60 | RL | 74.0 | 13830 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | 0 | |
| 1464 | 60 | RL | 78.0 | 9978 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | 0 | |
| 1465 | 120 | RL | 43.0 | 5005 | Pave | NaN | IR1 | HLS | AllPub | Inside | ... | 144 | 0 | |

5 rows × 79 columns



In [47]:

```
# basic cleaning (no)
repo.basic_cleaning(clean=False)
df = repo.get_data("basic")
print(df.shape)
df.head()
```

(1459, 79)

Out[47]:

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | LandContour | Utilities | LotConfig | ... | ScreenPorch | PoolArea | F |
|------|------------|----------|-------------|---------|--------|-------|----------|-------------|-----------|-----------|-----|-------------|----------|---|
| Id | | | | | | | | | | | | | | |
| 1461 | 20 | RH | 80.0 | 11622 | Pave | NaN | Reg | Lvl | AllPub | Inside | ... | 120 | 0 | |
| 1462 | 20 | RL | 81.0 | 14267 | Pave | NaN | IR1 | Lvl | AllPub | Corner | ... | 0 | 0 | |
| 1463 | 60 | RL | 74.0 | 13830 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | 0 | |
| 1464 | 60 | RL | 78.0 | 9978 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | 0 | |
| 1465 | 120 | RL | 43.0 | 5005 | Pave | NaN | IR1 | HLS | AllPub | Inside | ... | 144 | 0 | |

5 rows × 79 columns



In [48]:

```
# feature selction (n0)
repo.feature_selection(variance_selector=False)
```

```
df = repo.get_data("selected")
print(df.shape)
df.head()
```

(1459, 79)

Out[48]:

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | LandContour | Utilities | LotConfig | ... | ScreenPorch | PoolArea | F |
|------|------------|----------|-------------|---------|--------|-------|----------|-------------|-----------|-----------|-----|-------------|----------|---|
| Id | | | | | | | | | | | | | | |
| 1461 | 20 | RH | 80.0 | 11622 | Pave | NaN | Reg | Lvl | AllPub | Inside | ... | 120 | 0 | |
| 1462 | 20 | RL | 81.0 | 14267 | Pave | NaN | IR1 | Lvl | AllPub | Corner | ... | 0 | 0 | |
| 1463 | 60 | RL | 74.0 | 13830 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | 0 | |
| 1464 | 60 | RL | 78.0 | 9978 | Pave | NaN | IR1 | Lvl | AllPub | Inside | ... | 0 | 0 | |
| 1465 | 120 | RL | 43.0 | 5005 | Pave | NaN | IR1 | HLS | AllPub | Inside | ... | 144 | 0 | |

5 rows × 79 columns



In [49]:

```
# feature engineering
repo.feature_engineering(sub_class=sub_class)
df = repo.get_data("engineered")
print(df.shape)
df.head()
```

(1459, 85)

Out[49]:

| | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | LandContour | Utilities | LotConfig | ... | MoSold | YrSold | SaleTy |
|------|--|----------|-------------|---------|--------|-------|----------|-------------|-----------|-----------|------------|--------|--------|--------|
| Id | | | | | | | | | | | | | | |
| 1461 | 1-STORY 1946 & NEWER ALL STYLES | RH | 80.0 | 11622 | Pave | NaN | Reg | | Lvl | AllPub | Inside ... | 6 | 2010 | W |
| 1462 | 1-STORY 1946 & NEWER ALL STYLES | RL | 81.0 | 14267 | Pave | NaN | IR1 | | Lvl | AllPub | Corner ... | 6 | 2010 | W |
| 1463 | 2-STORY 1946 & NEWER | RL | 74.0 | 13830 | Pave | NaN | IR1 | | Lvl | AllPub | Inside ... | 3 | 2010 | W |
| 1464 | 2-STORY 1946 & NEWER | RL | 78.0 | 9978 | Pave | NaN | IR1 | | Lvl | AllPub | Inside ... | 6 | 2010 | W |
| 1465 | 1-STORY PUD (Planned Unit Development) - 1946 ... | RL | 43.0 | 5005 | Pave | NaN | IR1 | | HLS | AllPub | Inside ... | 1 | 2010 | W |

5 rows × 85 columns



In [50]:

```
# final mapping
df_test = df[X_train.columns]
print(df_test.shape)
df_test.head()
```

(1459, 72)

Out[50]:

| | MSSubClass | LotFrontage | LotArea | OverallQual | OverallCond | YearBuilt | YearRemodAdd | MasVnrArea | BsmtFinSF1 | BsmtFinSF2 | ... | Ga |
|------|--|-------------|---------|-------------|-------------|-----------|--------------|------------|------------|------------|-----|----|
| Id | | | | | | | | | | | | |
| 1461 | 1-STORY 1946 & NEWER ALL STYLES | 80.0 | 11622 | 5 | 6 | 1961 | 1961 | 0.0 | 468.0 | 144.0 | ... | |
| 1462 | 1-STORY 1946 & NEWER ALL STYLES | 81.0 | 14267 | 6 | 6 | 1958 | 1958 | 108.0 | 923.0 | 0.0 | ... | |
| 1463 | 2-STORY 1946 & NEWER | 74.0 | 13830 | 5 | 5 | 1997 | 1998 | 0.0 | 791.0 | 0.0 | ... | |
| 1464 | 2-STORY 1946 & NEWER | 78.0 | 9978 | 6 | 6 | 1998 | 1998 | 20.0 | 602.0 | 0.0 | ... | |
| 1465 | 1-STORY PUD (Planned Unit Development) - 1946 ... | 43.0 | 5005 | 8 | 5 | 1992 | 1992 | 0.0 | 263.0 | 0.0 | ... | |

5 rows × 72 columns



In [51]:

```
from Training import TestPredictor
# Get the predictions
tp = TestPredictor(test_data=df_test, model=linear_model)
print(type(tp))
tp
```

<class 'Training.TestPredictor'>

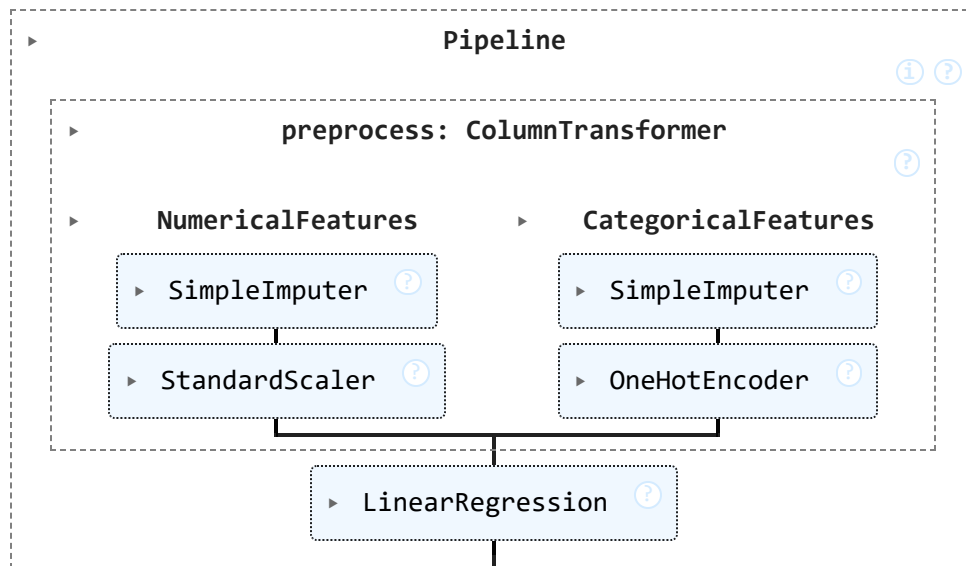
Out[51]: TestMapper on C:\Users\MY PC\Desktop\Projects\Regression\AmosHousePriceModelling

In [52]:

```
# Linear regression
linear_model = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("linear_model", LinearRegression())
    ]
)
```

```
# Training the model
linear_model.fit(X_train, y_train)
```

Out[52]:



```
In [53]: tp.predict()
pred = tp.get_data("prediction")
print(type(pred))
pred[:4]
```

```
<class 'numpy.ndarray'>
```

```
Out[53]: array([111934.25406071, 158050.42262291, 179127.82260186, 195072.54704813])
```

```
In [54]: tp.id_mapper(label="linear_regression")
sub = tp.get_data("mapped")
print(type(sub))
print(sub.shape)
sub.head()
```

```
<class 'pandas.core.frame.DataFrame'>
(1459, 1)
```

Out[54]:

SalePrice

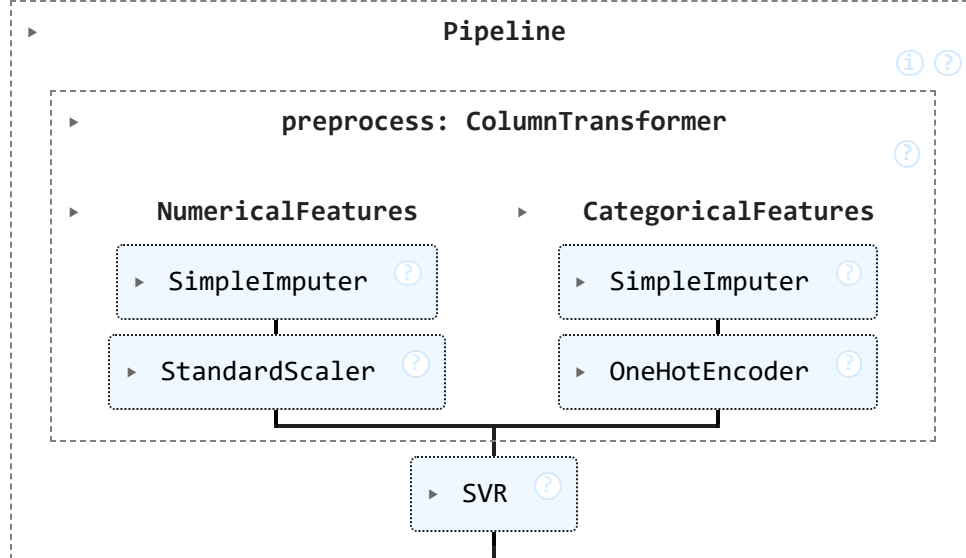
| Id | |
|-------------|---------------|
| 1461 | 111934.254061 |
| 1462 | 158050.422623 |
| 1463 | 179127.822602 |
| 1464 | 195072.547048 |
| 1465 | 185396.000891 |

9. Polynomial Features (degree 2) and SVR

To improve our model performance we need to include a polynomial of degree 2 to our support vector regressor (SVR)

```
In [55]: # building the pipeline
svr_pipeline = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("svr_model", SVR(kernel="poly", degree=2))
    ]
)
# Training the model
svr_model = svr_pipeline.fit(X_train, y_train)
svr_model
```

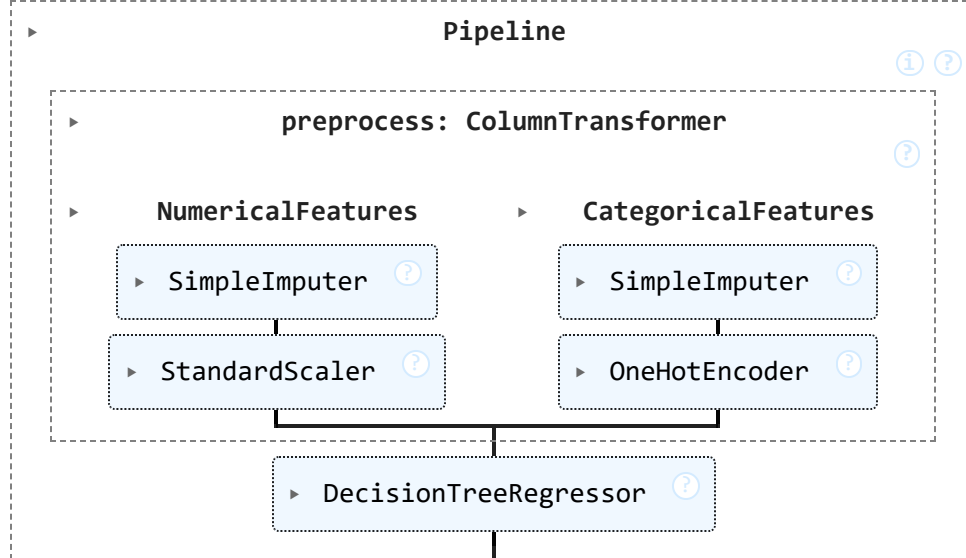
Out[55]:



10. Decision Tree Regressor

```
In [56]: # Building the pipeline with decision tree regressor
tree_pipeline = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("tree_model", DecisionTreeRegressor(random_state=42))
    ]
)
# Training the model
tree_model = tree_pipeline.fit(X_train, y_train)
tree_model
```


Out[56]:



We have now trained a decision tree model, next up is to evaluate our model and continue to do hyperparameter tuning.

```
In [57]: # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = tree_model.predict(X_train)
# With training data
mae = mean_absolute_error(y_train, pred)
mape = mean_absolute_percentage_error(y_train, pred)
cod = r2_score(y_train, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$0.0

Training MAPE: 0.0%

Training R2: 100.0%

```
In [58]: # Evaluation using validation set
# model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = tree_model.predict(X_val)
# With training data
mae = mean_absolute_error(y_val, pred)
mape = mean_absolute_percentage_error(y_val, pred)
cod = r2_score(y_val, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$26578.59
Training MAPE: 14.000000000000002%
Training R2: 56.99999999999999%

Definitely we have some element of overfitting, let's visualize this using a plot. We are going to use the LearningCurve class that has all the three modules.

```
In [59]: # get the learning curve class
lc = LearningCurve(estimator=tree_model, X=X, y=y)
print(type(lc))
lc
```

```
<class 'Training.LearningCurve'>
```

```
Out[59]: LearningCurve: <class 'sklearn.pipeline.Pipeline'>
```

```
In [60]: # Het Learning curve List data
lc.learning_curve()
tree_lc = lc.get_data("lc")
print(type(tree_lc))
tree_lc[:4]
```

```
[learning_curve] Training set sizes: [ 93 304 514 725 936]
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 out of 25 | elapsed: 3.0s finished
```

```
<class 'list'>
```

```
Out[60]: [array([ 93, 304, 514, 725, 936]),
          array([1., 1., 1., 1., 1.]),
          array([0.11831156, 0.54036325, 0.72716731, 0.69733103, 0.75523564])]
```

```
In [61]: # make the learning curve dataframe
lc.make_dataframe()
tree_lc_df = lc.get_data("df_lc")
print(type(tree_lc_df))
print(tree_lc_df.shape)
tree_lc_df
```

```
<class 'pandas.core.frame.DataFrame'>
(5, 3)
```

Out[61]:

| | Train Size | Train R2 | Validation R2 |
|----------|------------|----------|---------------|
| 0 | 93 | 1.0 | 0.118312 |
| 1 | 304 | 1.0 | 0.540363 |
| 2 | 514 | 1.0 | 0.727167 |
| 3 | 725 | 1.0 | 0.697331 |
| 4 | 936 | 1.0 | 0.755236 |

In [62]:

```
# Melting the dataframe
lc.melt_dataframe()
tree_lc_melt = lc.get_data("melt_lc")
print(type(tree_lc_melt))
print(tree_lc_melt.shape)
tree_lc_melt.head()
```

```
<class 'pandas.core.frame.DataFrame'>
(10, 3)
```

Out[62]:

| | Train Size | Set | R2 |
|----------|------------|----------|-----|
| 0 | 93 | Train R2 | 1.0 |
| 1 | 304 | Train R2 | 1.0 |
| 2 | 514 | Train R2 | 1.0 |
| 3 | 725 | Train R2 | 1.0 |
| 4 | 936 | Train R2 | 1.0 |

In [63]:

```
# making the plot
lc.plot_lc()
fig = lc.get_data()
print(type(fig))
save_plot(fname="decion_tree_model_learning_curve", filetype="plt", fig=fig)
fig.show()
```

```
<class 'plotly.graph_objs._figure.Figure'>
```

<Figure size 640x480 with 0 Axes>

As seen we are overfitting and we need to improve our model by doing the following:

1. Reduce model complexity- reducing the depth of the model
2. Do a cross validation (kfold)
3. hyperparameter tuning

But before we do that, let us get the maxim depth of our model.

```
In [64]: # trianing different trees
train_acc = []
val_acc = []
d_params = range(1, 10)
for d in d_params:
```

```

# training the model
model = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("model", DecisionTreeRegressor(max_depth=d, random_state=42))
    ]
)
model.fit(X_train, y_train)
# R2
train_acc.append(r2_score(y_train, model.predict(X_train)))
val_acc.append(r2_score(y_val, model.predict(X_val)))

```

```

In [65]: # building a dataframe
df_result = pd.DataFrame(
    {
        "Depth": d_params,
        "Train Accuracy": train_acc,
        "Validation Accuracy": val_acc
    }
)
df_result

```

```

Out[65]:

```

| | Depth | Train Accuracy | Validation Accuracy |
|---|-------|----------------|---------------------|
| 0 | 1 | 0.453900 | 0.334842 |
| 1 | 2 | 0.629325 | 0.501178 |
| 2 | 3 | 0.762974 | 0.628356 |
| 3 | 4 | 0.830680 | 0.702260 |
| 4 | 5 | 0.886568 | 0.636079 |
| 5 | 6 | 0.924428 | 0.597267 |
| 6 | 7 | 0.951208 | 0.762117 |
| 7 | 8 | 0.971141 | 0.814099 |
| 8 | 9 | 0.982418 | 0.623959 |

```

In [66]: # melting the dataframe,
result_melt = pd.melt(
    frame=df_result,
    id_vars="Depth",
    value_vars=["Train Accuracy", "Validation Accuracy"],
    value_name="Accuracy",

```

```
    var_name="Set"  
)  
result_melt.head()
```

Out[66]:

| | Depth | Set | Accuracy |
|---|-------|----------------|----------|
| 0 | 1 | Train Accuracy | 0.453900 |
| 1 | 2 | Train Accuracy | 0.629325 |
| 2 | 3 | Train Accuracy | 0.762974 |
| 3 | 4 | Train Accuracy | 0.830680 |
| 4 | 5 | Train Accuracy | 0.886568 |

In [67]:

```
# plotting  
fig = px.line(  
    data_frame=result_melt,  
    x = "Depth",  
    y = "Accuracy",  
    color = "Set",  
    title="Decision Tree Model: Training and validation accuracy curves"  
)  
fig.update_layout(  
    xaxis_title="Depth of the Tree",  
    yaxis_title="Accuracy",  
    legend_title="Accuracy"  
)  
fig.show()
```

As seen in the plot above a depth of seven is having the highest validation accuracy. Let us train our final model using the depth of seven first and define other hyperparameters that will help reduce overfitting.

```
In [68]: params = {  
    "tree_model__min_samples_split": [3,4,5,6],  
    "tree_model__min_samples_leaf": [1,2,3,4]  
}  
params
```

```
Out[68]: {'tree_model__min_samples_split': [3, 4, 5, 6],  
          'tree_model__min_samples_leaf': [1, 2, 3, 4]}
```

```
In [69]: # model pipeline  
tree_model_pipeline = Pipeline(  
    [  
        
```

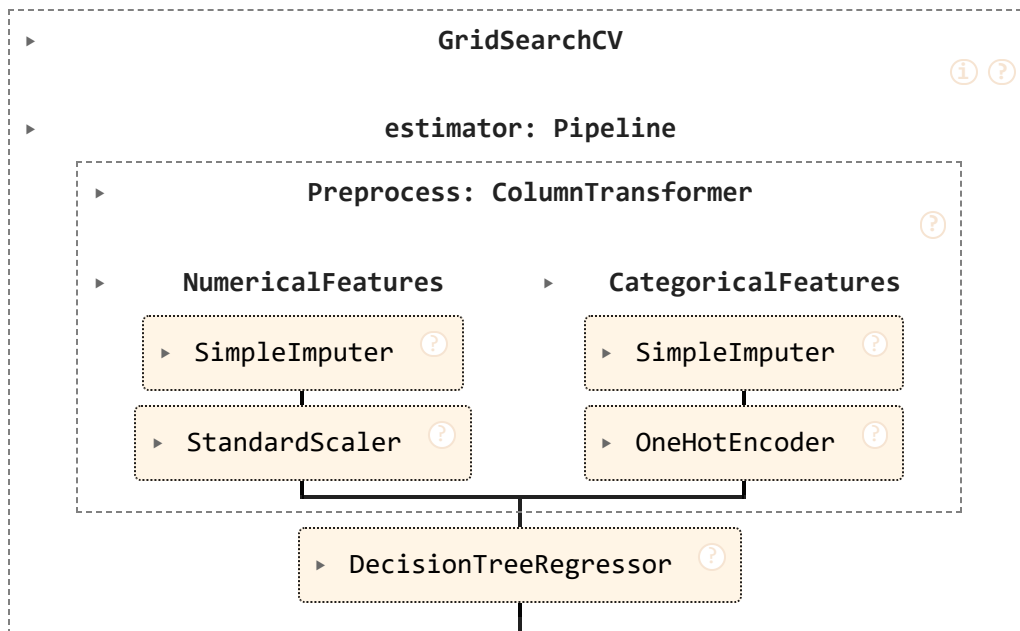
```

        ("Preprocess", col_pipeline),
        ("tree_model", DecisionTreeRegressor(max_depth=7, random_state=42))
    ]
)

# Grid search
tree_model_cv = GridSearchCV(
    estimator=tree_model_pipeline,
    param_grid=params,
    n_jobs=-1,
    cv=5,
    verbose=1
)
tree_model_cv

```

Out[69]:



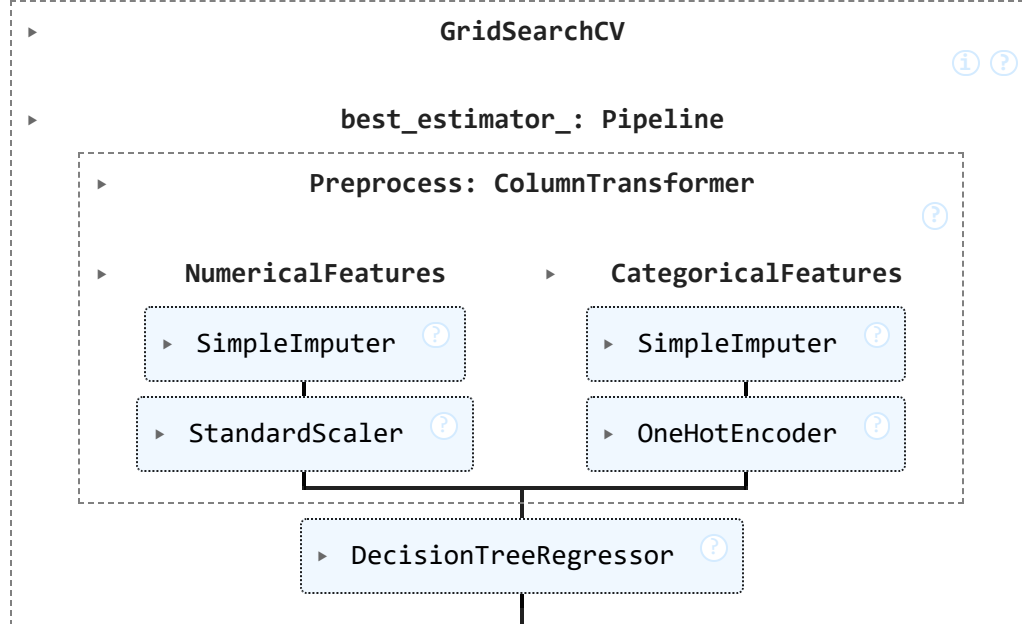
```

In [70]: # fitting the model
tree_model_cv.fit(X_train, y_train)

```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

Out[70]:



```
In [71]: # getting the best paramers
tree_model_cv.best_params_
```

```
Out[71]: {'tree_model__min_samples_leaf': 1, 'tree_model__min_samples_split': 4}
```

```
In [72]: # Betting the model cv results into a dataframe
cv_result = pd.DataFrame.from_dict(tree_model_cv.cv_results_)
print(type(cv_result))
print(cv_result.shape)
cv_result.head()
```

```
<class 'pandas.core.frame.DataFrame'>
(16, 15)
```

Out[72]:

| | mean_fit_time | std_fit_time | mean_score_time | std_score_time | param_tree_model__min_samples_leaf | param_tree_model__min_samples_split | |
|---|---------------|--------------|-----------------|----------------|------------------------------------|-------------------------------------|------|
| 0 | 0.209090 | 0.027543 | 0.058873 | 0.017155 | 1 | 3 | {'tr |
| 1 | 0.182238 | 0.010539 | 0.039121 | 0.003061 | 1 | 4 | {'tr |
| 2 | 0.174386 | 0.005411 | 0.045016 | 0.005904 | 1 | 5 | {'tr |
| 3 | 0.170947 | 0.010402 | 0.042823 | 0.006001 | 1 | 6 | {'tr |
| 4 | 0.226941 | 0.033243 | 0.060980 | 0.012052 | 2 | 3 | {'tr |

In [73]:

```
# making splits score dataframe
score_col = ["rank_test_score", "split0_test_score", "split1_test_score", "split2_test_score", "split3_test_score", "split4_test_s
split_score = cv_result[score_col]
# melting the dataframe
split_score = pd.melt(frame=split_score, var_name="Set", value_name="Score", id_vars="rank_test_score")
split_score.head()
```

Out[73]:

| | rank_test_score | Set | Score |
|---|-----------------|-------------------|----------|
| 0 | 13 | split0_test_score | 0.772907 |
| 1 | 1 | split0_test_score | 0.776451 |
| 2 | 8 | split0_test_score | 0.759828 |
| 3 | 16 | split0_test_score | 0.762805 |
| 4 | 2 | split0_test_score | 0.750379 |

In [74]:

```
# plotting
fig = px.line(
    data_frame=split_score,
    x=split_score.index,
    y="Score",
    color="Set",
    title="Hyperameter Sets: Decision Tree Model"
)
fig.update_layout(
    xaxis_title="Training Index",
```

```
    yaxis_title="Accuracy (%)",  
    legend_title="Split Type"  
)  
fig.show()
```

As seen some model had really high accuracies but now let us get the best model and parameters.

```
In [75]: tree_model_cv.best_params_
```

```
Out[75]: {'tree_model__min_samples_leaf': 1, 'tree_model__min_samples_split': 4}
```

Now that we have our best model, let us evaluate it.

```
In [76]: # model evaluation using mae, mape, coefficient of difference(cod, R2)  
pred = tree_model_cv.predict(X_train)
```

```
# With training data
mae = mean_absolute_error(y_train, pred)
mape = mean_absolute_percentage_error(y_train, pred)
cod = r2_score(y_train, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$11483.79

Training MAPE: 7.000000000000001%

Training R2: 95.0%

```
In [77]: # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = tree_model_cv.predict(X_val)
# With training data
mae = mean_absolute_error(y_val, pred)
mape = mean_absolute_percentage_error(y_val, pred)
cod = r2_score(y_val, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$22159.47

Training MAPE: 12.0%

Training R2: 73.0%

We have tried to reduce overfitting but one more thing to be done is increasing the data. Since this is a kaggle competition, then we will not do that. we will investigate this property using a learning curve.

```
In [78]: lc = LearningCurve(estimator=tree_model_cv, X=X, y=y)
print(type(lc))
lc
```

<class 'Training.LearningCurve'>

Out[78]: LearningCurve: <class 'sklearn.model_selection._search.GridSearchCV'>

```
In [79]: # Learning curve list
lc.learning_curve()
ls = lc.get_data("lc")
print(type(ls))
ls[:4]
```

[learning_curve] Training set sizes: [93 304 514 725 936]

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 25 out of 25 | elapsed: 1.4min finished

<class 'list'>

```
Out[79]: [array([ 93, 304, 514, 725, 936]),
          array([0.91356261, 0.94798811, 0.95389146, 0.94835213, 0.9373437 ]),
          array([0.27786857, 0.65022828, 0.72233115, 0.75555916, 0.77462749])]
```

```
In [80]: # getting the data
lc.make_dataframe()
ld = lc.get_data("df_lc")
print(type(ld))
print(ld.shape)
ld.head()
```

```
<class 'pandas.core.frame.DataFrame'>
(5, 3)
```

```
Out[80]:
```

| | Train Size | Train R2 | Validation R2 |
|---|------------|----------|---------------|
| 0 | 93 | 0.913563 | 0.277869 |
| 1 | 304 | 0.947988 | 0.650228 |
| 2 | 514 | 0.953891 | 0.722331 |
| 3 | 725 | 0.948352 | 0.755559 |
| 4 | 936 | 0.937344 | 0.774627 |

```
In [81]: # getting the melted data
lc.melt_dataframe()
ld = lc.get_data("melt_lc")
print(type(ld))
print(ld.shape)
ld.head()
```

```
<class 'pandas.core.frame.DataFrame'>
(10, 3)
```

```
Out[81]:
```

| | Train Size | Set | R2 |
|---|------------|----------|----------|
| 0 | 93 | Train R2 | 0.913563 |
| 1 | 304 | Train R2 | 0.947988 |
| 2 | 514 | Train R2 | 0.953891 |
| 3 | 725 | Train R2 | 0.948352 |
| 4 | 936 | Train R2 | 0.937344 |

```
In [82]: # getting the data
lc.plot_lc()
fig = lc.get_data()
print(type(fig))
fig.show()
```

```
<class 'plotly.graph_objs._figure.Figure'>
```

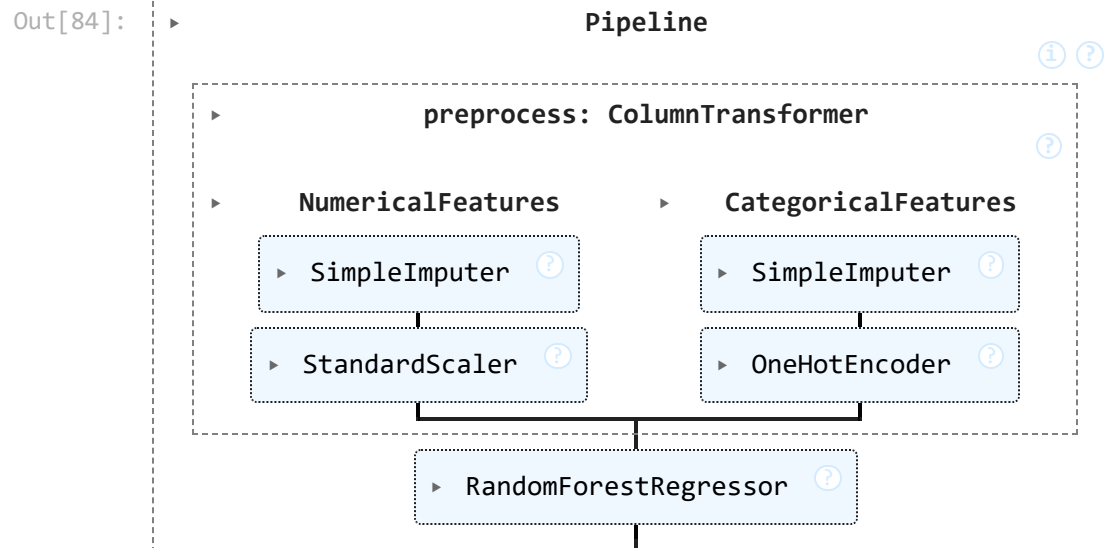
Definitely this plot confirms that we need more data which we don't have maybe we can improve our model using another method. One thing we have done successfully is that we have been able to reduce overfitting as seen. We will start doing bagging and boosting learning methods to try and see if our model performance will increase and we reduce overfitting.

```
In [83]: # saving the model
save_model(mname="decision_tree_model", model=tree_model_cv)
```

11. Bagging Model

Random forest model is one of the bagging models. We are going to train that and see the progress.

```
In [84]: # Model pipeline
forest_model = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("forest_model", RandomForestRegressor(max_depth=7))
    ]
)
# Training the model
forest_model.fit(X_train, y_train)
```



```
In [85]: # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = forest_model.predict(X_train)
# With training data
mae = mean_absolute_error(y_train, pred)
mape = mean_absolute_percentage_error(y_train, pred)
cod = r2_score(y_train, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$10335.71

Training MAPE: 6.0%

Training R2: 96.0%

```
In [86]: # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = forest_model.predict(X_val)
# With training data
mae = mean_absolute_error(y_val, pred)
mape = mean_absolute_percentage_error(y_val, pred)
cod = r2_score(y_val, pred)
print(f"Validation MAE: ${np.round(mae, 2)}")
print(f"Validation MAPE: {100 * np.round(mape, 2)}%")
print(f"Validation R2: {100 * np.round(cod, 2)}%")
```

Validation MAE: \$16426.49

Validation MAPE: 9.0%

Validation R2: 89.0%

Actually this model is performing better. We will first check for overfitting and then do a hyperparameter tuning.

```
In [87]: lcf = LearningCurve(estimator=forest_model, X=X, y=y)
print(type(lcf))
lcf
```

<class 'Training.LearningCurve'>

Out[87]: LearningCurve: <class 'sklearn.pipeline.Pipeline'>

```
In [88]: # Getting the Learning curve list
lcf.learning_curve()
lcf_list = lcf.get_data("lc")
print(type(lcf_list))
lcf
```

[learning_curve] Training set sizes: [93 304 514 725 936]

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 25 out of 25 | elapsed: 32.1s finished

<class 'list'>

Out[88]: LearningCurve: <class 'sklearn.pipeline.Pipeline'>

```
In [89]: # getting the dataframe
lcf.make_dataframe()
lcf_df = lcf.get_data("df_lc")
print(type(lcf_df))
print(lcf_df.shape)
lcf_df
```

<class 'pandas.core.frame.DataFrame'>

(5, 3)

Out[89]:

| | Train Size | Train R2 | Validation R2 |
|----------|------------|----------|---------------|
| 0 | 93 | 0.952477 | 0.798010 |
| 1 | 304 | 0.971097 | 0.831800 |
| 2 | 514 | 0.967815 | 0.850644 |
| 3 | 725 | 0.964742 | 0.855944 |
| 4 | 936 | 0.962920 | 0.872585 |

In [90]:

```
# melting the dataframe
lcf.melt_dataframe()
lcf_melt = lcf.get_data("melt_lc")
print(type(lcf_melt))
print(lcf_melt.shape)
lcf_melt.head()
```

<class 'pandas.core.frame.DataFrame'>
(10, 3)

Out[90]:

| | Train Size | Set | R2 |
|----------|------------|----------|----------|
| 0 | 93 | Train R2 | 0.952477 |
| 1 | 304 | Train R2 | 0.971097 |
| 2 | 514 | Train R2 | 0.967815 |
| 3 | 725 | Train R2 | 0.964742 |
| 4 | 936 | Train R2 | 0.962920 |

In [91]:

```
# plotting the learning curve
lcf.plot_lc()
fig = lcf.get_data()
print(type(fig))
fig.show()
```

<class 'plotly.graph_objs._figure.Figure'>

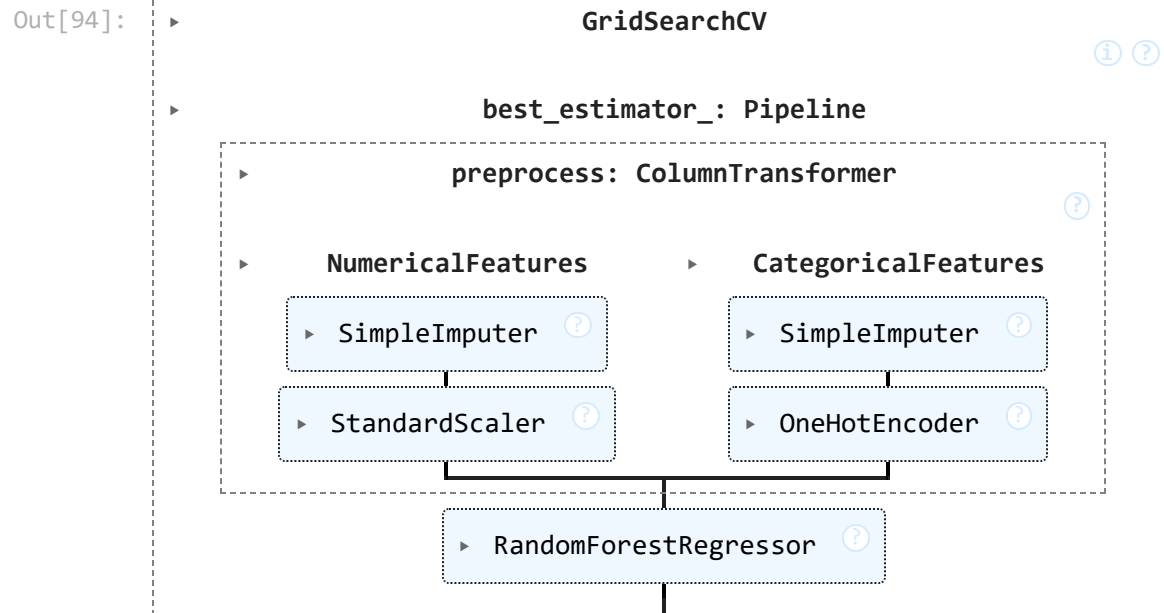
With more data we can actually get better results. Our model is not overfitting but we will try different hyperparameter to see if we will get a better optimal parameters

```
In [93]: f_params = {  
    "forest_model__n_estimators": range(20, 100, 20),  
    "forest_model__max_depth": range(4,12,2),  
    "forest_model__min_samples_split": [4,6,8,10],  
    "forest_model__min_samples_leaf": [1,2,3,4]  
}  
f_params
```

```
Out[93]: {'forest_model__n_estimators': range(20, 100, 20),
          'forest_model__max_depth': range(4, 12, 2),
          'forest_model__min_samples_split': [4, 6, 8, 10],
          'forest_model__min_samples_leaf': [1, 2, 3, 4]}
```

```
In [94]: # model pipeline
forest_model_cv = Pipeline(
    [
        ("preprocess", col_pipeline),
        ("forest_model", RandomForestRegressor(random_state=42))
    ]
)
# Cross validation
forest_model_cv = GridSearchCV(
    estimator=forest_model_cv,
    param_grid=f_params,
    cv=5,
    n_jobs=-1,
    verbose=1
)
# Training the model
forest_model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 256 candidates, totalling 1280 fits



```
In [99]: # Getting the model best parameters
forest_model_cv.best_params_
```


```
Out[99]: {'forest_model__max_depth': 10,  
         'forest_model__min_samples_leaf': 2,  
         'forest_model__min_samples_split': 4,  
         'forest_model__n_estimators': 60}
```

```
In [100... # Getting the cv results  
# Betting the model cv results into a dataframe  
cv_result = pd.DataFrame.from_dict(forest_model_cv.cv_results_)  
print(type(cv_result))  
print(cv_result.shape)  
cv_result.head()
```

```
<class 'pandas.core.frame.DataFrame'>  
(256, 17)
```

Out[100...

| | mean_fit_time | std_fit_time | mean_score_time | std_score_time | param_forest_model__max_depth | param_forest_model__min_samples_leaf | param |
|---|---------------|--------------|-----------------|----------------|-------------------------------|--------------------------------------|-------|
| 0 | 0.639304 | 0.020015 | 0.053631 | 0.008212 | 4 | 1 | |
| 1 | 1.131988 | 0.033393 | 0.059633 | 0.011264 | 4 | 1 | |
| 2 | 1.886548 | 0.151949 | 0.072804 | 0.018020 | 4 | 1 | |
| 3 | 2.473378 | 0.201929 | 0.067602 | 0.002350 | 4 | 1 | |
| 4 | 0.597211 | 0.006840 | 0.046773 | 0.001308 | 4 | 1 | |

◀  ▶

```
In [101... # making splits score dataframe  
score_col = ["rank_test_score", "split0_test_score", "split1_test_score", "split2_test_score", "split3_test_score", "split4_test_s  
split_score = cv_result[score_col]  
# melting the dataframe  
split_score = pd.melt(frame=split_score, var_name="Set", value_name="Score", id_vars="rank_test_score")  
split_score.head()
```

Out[101...

| | rank_test_score | Set | Score |
|---|-----------------|-------------------|----------|
| 0 | 249 | split0_test_score | 0.832552 |
| 1 | 234 | split0_test_score | 0.837057 |
| 2 | 221 | split0_test_score | 0.842257 |
| 3 | 211 | split0_test_score | 0.841496 |
| 4 | 250 | split0_test_score | 0.832640 |

In [102...

```
# plotting
fig = px.line(
    data_frame=split_score,
    x=split_score.index,
    y="Score",
    color="Set",
    title="Hyperparameter Sets: Decision Tree Model"
)
fig.update_layout(
    xaxis_title="Training Index",
    yaxis_title="Accuracy (%)",
    legend_title="Split Type"
)
fig.show()
```

With that information let us evaluation our best model.

```
In [107... # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = forest_model_cv.best_estimator_.predict(X_train)
# With training data
mae = mean_absolute_error(y_train, pred)
mape = mean_absolute_percentage_error(y_train, pred)
cod = r2_score(y_train, pred)
print(f"Training MAE: ${np.round(mae, 2)}")
print(f"Training MAPE: {100 * np.round(mape, 2)}%")
print(f"Training R2: {100 * np.round(cod, 2)}%")
```

Training MAE: \$7489.56

Training MAPE: 5.0%

Training R2: 97.0%

```
In [108... # model evaluation using mae, mape, coefficient of difference(cod, R2)
pred = forest_model_cv.best_estimator_.predict(X_val)
# With training data
mae = mean_absolute_error(y_val, pred)
mape = mean_absolute_percentage_error(y_val, pred)
cod = r2_score(y_val, pred)
print(f"Validation MAE: ${np.round(mae, 2)}")
print(f"Validation MAPE: {100 * np.round(mape, 2)}%")
print(f"Validation R2: {100 * np.round(cod, 2)}%")
```

Validation MAE: \$16474.68

Validation MAPE: 9.0%

Validation R2: 88.0%

This model is performing pretty well.

Let us save map our ids for submission and finally save the model.

```
In [110... # id mapping
tcf = TestPredictor(test_data=df_test, model=forest_model_cv)
print(type(tcf))
tcf
```

<class 'Training.TestPredictor'>

Out[110... TestMapper on C:\Users\MY PC\Desktop\Projects\Regression\AmosHousePriceModelling

```
In [112... # Getting the predictions
tcf.predict()
pred_f = tcf.get_data("prediction")
print(type(pred_f))
pred_f[:4]
```

<class 'numpy.ndarray'>

Out[112... array([126023.64599443, 156174.380588 , 190244.77318438, 179339.0265955])

```
In [114... # mapping the ids
tcf.id_mapper(label="forest")
sub_f = tcf.get_data("mapped")
print(type(df))
print(sub_f.shape)
sub_f.head()
```

<class 'pandas.core.frame.DataFrame'>
(1459, 1)

Out[114...

SalePrice

| | Id |
|-------------|---------------|
| 1461 | 126023.645994 |
| 1462 | 156174.380588 |
| 1463 | 190244.773184 |
| 1464 | 179339.026595 |
| 1465 | 193461.518541 |

In [116...

```
val_pred_f = forest_model_cv.predict(X_val)
```

In [117...

```
# Scatter plot  
fig = scatter_plot(x=val_X_t, y=y_val, y_pred=val_pred_f, label="Random Forest Model")  
save_plot(fname="LinerModelScatter", filetype="plt")  
fig.show()
```


<Figure size 640x480 with 0 Axes>

```
In [118... # saving the model  
save_model(mname="RandomFores_model", model=forest_model_cv)
```

12. Boosting Model

We will try a boosting model, gradient boosting model and see it's performance. If it is going to perform better than the rest of the models.

```
In [ ]:
```